

EDUCATION

---

**University of California, Los Angeles**

*Bachelor of Science, Computer Science and Mathematics*

- **GPA: 3.93/4.00** (Dean’s honor list, 4 Quarters)
- **Extracurriculars:** William F. Sharpe Fellow, Bruin AI, Bruin Quant Traders, Founder- Olympic Robotics

**Expected: May 2028**

WORK EXPERIENCE

---

**BigML Lab @ UCLA** | *Undergraduate Researcher*

- Prof. Baharan Mirzasoleiman’s lab. Efficient GRPO and DPO algorithms for post-training LLMs.

**EXL** | *Data Management Intern (EXLData.ai)*

- Built AI agents to analyze and chunk internal repositories; developed a RAG-Fusion and KB auto-update CI/CD pipeline with LangChain and LangGraph; used by 40 employees and 100 clients; leading to 20% time saved.
- Resolved production failures across Databricks, Kubernetes, Helm, and Kafka; introduced pre-push secret scanning, Kubernetes admission controls, and CI security gates, reducing flagged security findings by 85%.

**Amazon Web Services (AWS)** | *AI Software Engineering Intern (contractual)*

- Collaborated to build production-grade RAG agent(website) using AWS AgentCore, to serve ~25k UCLA members.
- Implemented hybrid search (BM25 + semantic), Reciprocal Rank Fusion, context-injection, and user sessions; reduced latency from 38s to 5.5s via routing decisions, model selection, data caching and cron jobs for cold starts.

**Diro – Internet Original Documents** | *Software Engineering Intern (AI)*

- Worked with a team of 5 to build a production-grade AI usage tracking platform using LiteLLM, Langfuse, PostgreSQL and GCP, providing full-stack observability for company-wide AI usage, reducing costs by 28%.

Los Angeles, CA | **Aug 2026-Present**

New York, NY | **Summer 2026**

Los Angeles, CA | **Spring 2026**

New York, NY | **Summer 2025**

RESEARCH EXPERIENCE

---

**Asymmetric Collapse in Model Merging: When Refusal Overwrites Recognition- COLM 2026, AIW and Adv-ML**

- Evaluated Linear, SLERP, TIES, and DARE-TIES model merges of specialized Gemma-3-1B models; found that refusal transferred at 81–85% while recognition collapsed to at most 12.9%. Attributed to attention head changes.

**Predicting the Next State Is Not Enough: JEPA Representations for Lean Theorem Proving:** Submitted to EMNLP

- Applied JEPA (self-supervised learning) to Lean 4 prover, outperforming baseline by 5 percentage points.

PROJECTS

---

**Akatosh AI. [Website]**

*Deep Learning, REST APIs, Full-Stack Development, Vercel*

- Designed and trained a deep neural network to score leads; owned end-to-end architecture, APIs, integration expansion, and Vercel deployment with focus on performance, reliability, and scalability. 50+ users from Adobe.

**Cactus Compute (YC S25), Open Source Contributor. [GitHub]**

*C++, OpenMP, quantization, KV-cache quantization, tensor graph*

- Evaluated TurboQuant and KV-cache quantization strategies helping achieve a 7x memory reduction with minimal quality loss; contributed fixes and implementation proposals to the C++ inference engine.

**Clubapply.ai. [GitHub]**

*AWS (EC2, ECS, Lambda), Python, Docker, FastAPI, S3, IAM, RAG*

- Agentic platform with 8 orchestration agents for college club applications. Built multimodal mock interviewer. Winner, best use of AWS technology at AWS x BruinAI hackathon 2025.

**Dec 2025–May 2026**

**Mar 2026–Present**

**Oct 2025**